

실시간 시스템 개념 및 실시간성 보장 기법

Jinkyu Lee

Dept. of Computer Science and Engineering,
Sungkyunkwan University (SKKU)

Instructor Information

■ Jinkyu Lee 이진규

■ Associate Professor

- Department of Computer Science and Engineering 소프트웨어학과, SKKU (2014.3~)

■ Postdoctoral Researcher

- Department of EECS, University of Michigan
 - Advisor: Prof. Kang G. Shin

■ B.S., M.S., Ph.D.

- Department of Computer Science 전산학부, KAIST
 - Advisor: Prof. Insik Shin & Ikjun Yeom

■ Contact information

- Email. jinkyu.lee@skku.edu
- Website. <https://rtclskku.github.io/website/jinkyulee.html>
- Real-Time Computing Lab@SKKU. <https://rtclskku.github.io/website>


Instructor Information

■ Research areas

- Real-time system
 - ML (Machine learning) for RT, RT for ML
- Software-defined battery management
 - Battery management systems
- Mobile system

Specialty: **Real-time scheduling**

- System design and analysis with **timing guarantees**
- QoS support
- Resource management



Instructor Information

■ Achievements: research

■ IEEE **RTSS** (Real-Time Systems Symposium)

Top 1 RT conference, about 30 papers published per year

- 19 papers published

■ IEEE **RTAS** (Real-Time and Embedded Technology and Application Symposium)

Top 2 RT conference, about 30 papers published per year

- 6 papers published

■ SCI(E) Journals

- 53 papers, including 17 IEEE/ACM transactions/journals

■ Achievements: professional activities

■ TPC (Technical Program Committee) co-chair: IEEE RTCSA 2020

■ TPC: IEEE **RTSS** 2018, 2020, 2021 and 2022

■ TPC: IEEE **RTAS** 2020, 2022 and 2023

Instructor Information

- Research example
 - How to capture users' grip for smartphones?



<https://www.youtube.com/watch?v=FvQ87wmS6kk>

SmartGrip: Grip Sensing System for Commodity Mobile Devices Through Sound Signals
Personal and Ubiquitous Computing 2020

Announcement

- 카메라를 켜주세요.
- Zoom 참여이름을 본인의 이름으로 셋팅해주세요.
- 과제: 오늘 수업이 끝나기 전까지 개인별로 질문 한개씩을 해주세요.

- 출석체크

Overall Schedule

- **실시간 시스템 개념**: 실시간 시스템의 개념을 설명하고, 실시간성 보장에 관한 기본적인 배경 지식을 공부함.
 - 09:00 – 10:15 강의
 - 10:30 – 11:45 **Case Study**
- 점심식사
 - 12:00 – 13:00
- **실시간성 보장 기법**: 주어진 마감 시간 내에 작업 완료를 “보장” 할 수 있는 기법들에 대해 공부함.
 - 13:00 – 14:15 강의
 - 14:30 – 15:45 과제
- 강의 노트 링크: https://rtclskku.github.io/website/papers/20230110_SSH.pdf

Real-Time System

- Systems that operate with time constraints (deadlines)
 - Important to produce accurate results **before deadlines**
 - Usually for safety-critical systems or mission-critical systems


Automotive



Avionics



Medical Devices



Robots



Military



Smartphones



Multimedia



Factory Automation



Real-Time System

- Systems that operate with time constraints (deadlines)
 - Important to produce accurate results **before deadlines**
 - Usually for safety-critical systems or mission-critical systems

**This is your airbag
on time.**

From Honda

Real-Time Systems

■ Definition

- Systems whose correctness depend on their temporal aspects as well as their functional aspects

■ Performance measure

- Timeliness on timing constraints (deadlines)
- Speed/average case performance are less significant.

■ Key property

- Predictability on timing constraints

Misconceptions about Real-Time Systems

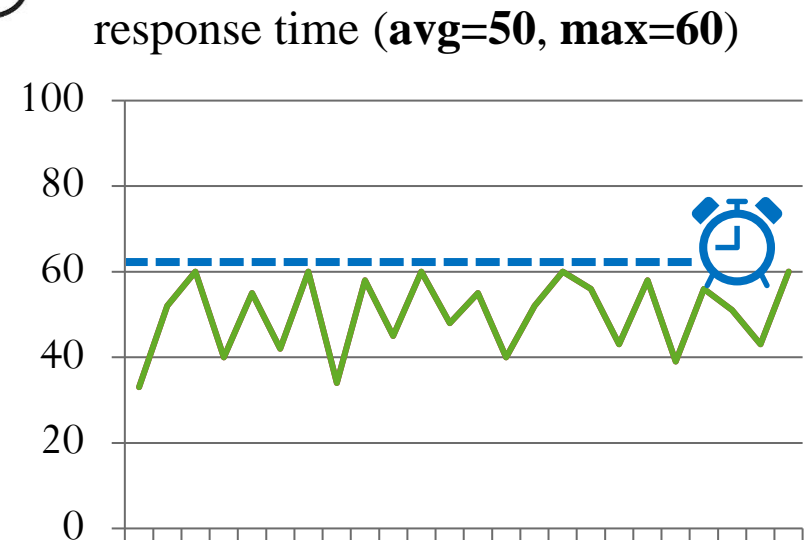
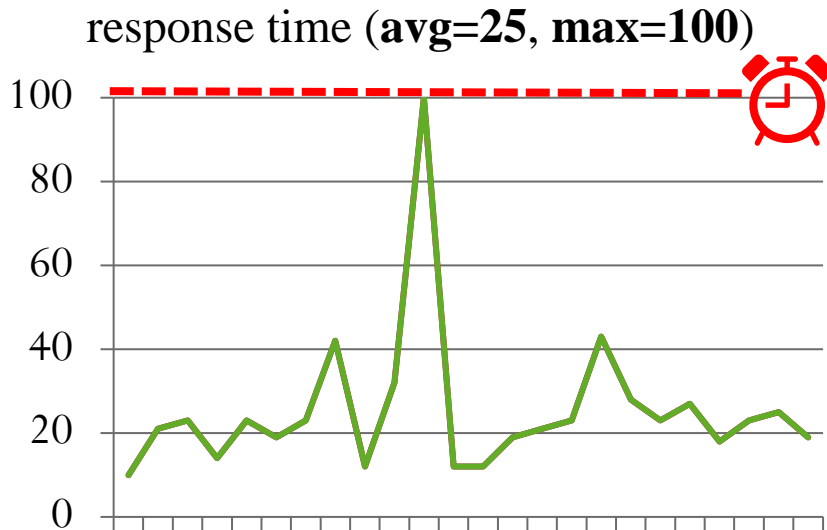
- Real-time \neq fast
 - Rather predictable than fast
 - “A man drowned in a river with an average depth of 20 centimeters



Real-Time System

- Systems that operate with time constraints (deadlines)
 - Important to produce accurate results **before deadlines**
 - Usually for safety-critical systems or mission-critical systems

VS



Real-Time Systems: Hard vs. Soft

■ Hard real-time systems

- Disastrous or very serious consequences may occur if a deadline is missed
- Validation is essential
 - Can all the deadlines be met, even under worst-case scenario?
- Deterministic guarantees

■ Soft real-time systems

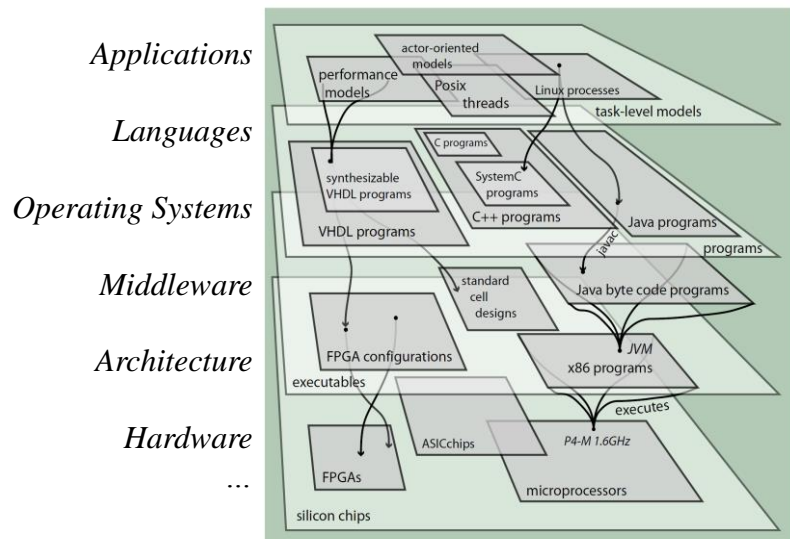
- Ideally, the deadline should be met for maximum performance. The performance degrades in case of deadline misses.
- Best effort approaches / statistical guarantees

Real-Time Systems: *Research Topics*

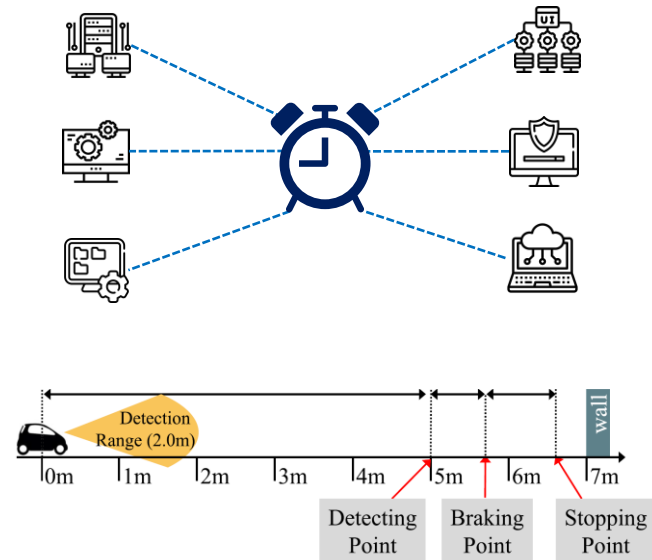
Design *time-predictable* computing systems

+

Timing guarantee and analysis of computing systems



Cyber Physical Systems: Design Challenges, IEEE ISORC, 2008.



MC-SDN: Supporting Mixed-Criticality Real-Time Communication Using Software-Defined Networking, IEEE IoT Journal, 2019.

Real-Time Systems: *Research Topics*

- Research example

- How to provide timing guarantees for software-defined networks?

MC-SDN

Case Study: 1/10 Scale Autonomous Car

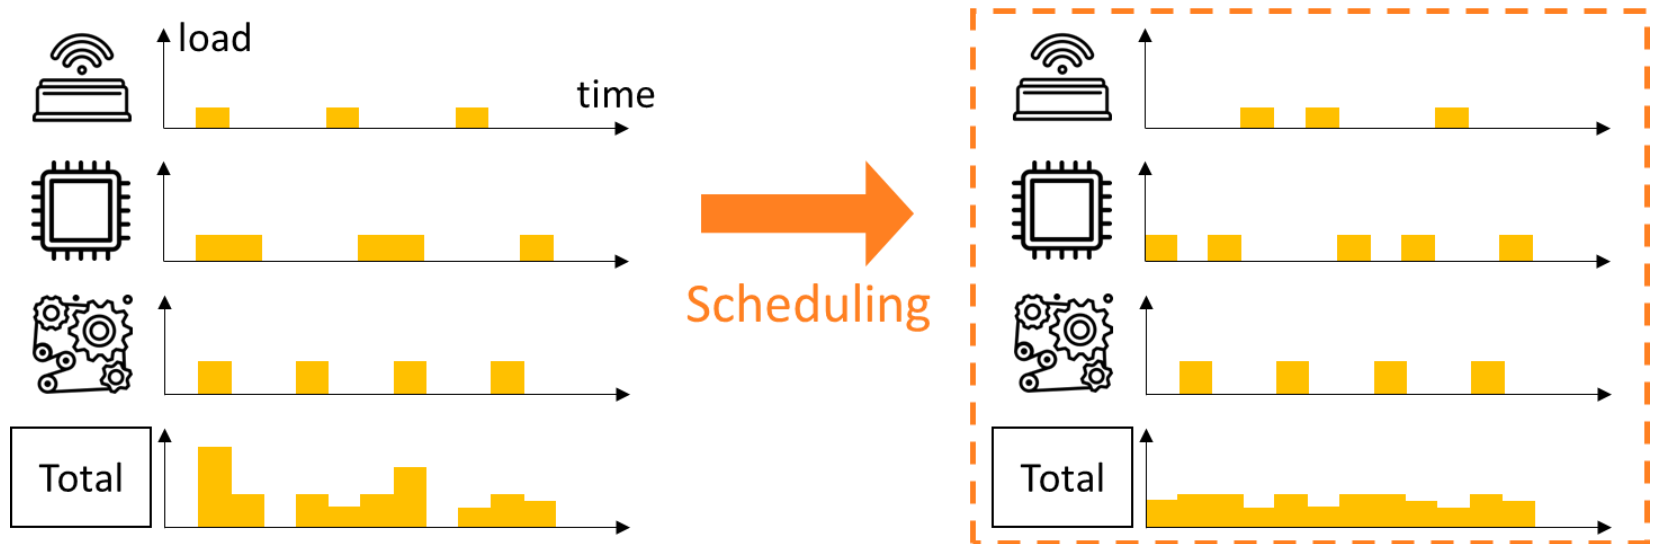
Kilho Lee, Taejune Park, Minsu Kim, Hoon Sung Chwa,
Jinkyu Lee, Seungwon Shin, and Insik Shin.

MC-SDN: Supporting Mixed-Criticality Real-Time Communication Using Software-Defined Networking
IEEE **RTSS** 2018

Real-Time Systems: *Research Topics*

■ Research example

- How to minimize battery aging while providing timing guarantees of real-time tasks?



Battery Aging Deceleration for Power-Consuming Real-Time Systems
IEEE **RTSS** 2019

Real-Time Systems: *Research Topics*

■ Research example

- How to fully utilize CPU *and* GPU for time-predictable execution for DNN tasks?

Timing analysis

$$I_k(w_k^*) = \begin{cases} \sum_{\tau_k \in \text{hp}(\tau_i)} \left(1 + \left\lfloor \frac{w_k^*}{T_k} \right\rfloor\right) \cdot (C_k^* + \sum_{\phi_{k,p}^* \in \phi_k^*} O(\phi_{k,p}^*)) & \text{if } \phi_k^* = \emptyset, \quad (5a) \\ \sum_{\tau_k \in \text{hp}(\tau_i)} \left(1 + \left\lfloor \frac{w_k^*}{T_k} \right\rfloor\right) \cdot (C_k^* + \sum_{\phi_{k,q}^* \in \phi_k^*} O(\phi_{k,q}^*)) & \text{if } \phi_k^* \neq \emptyset, \quad (5b) \\ (5a) + (5b), & \text{otherwise,} \quad (5c) \end{cases}$$

$$B_i = |\phi_i^*| \cdot B_{\max}^G(\tau_i) + |\phi_i^*| \cdot B_{\max}^C(\tau_i), \quad (6)$$

$$\text{where } B_{\max}^C(\tau_i) = \max_{j: \tau_j \in \text{hp}(\tau_i) \wedge \tau_{j,j} \in \Lambda_i^C} (C_{i,j}^C - 1), \quad (7)$$

$$B_{\max}^G(\tau_i) = \max_{k: \tau_m \in \text{hp}(\tau_i) \wedge \tau_{m,k} \in \Lambda_{i,m}^G} (C_{m,k}^G - 1), \quad (8)$$



Layer allocation

Algorithm 1 Layer-by-layer CPU/GPU allocation

```

1: Input:  $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$ 
2: Output:  $\Lambda = \{\{\Lambda_1^C, \Lambda_1^G\}, \{\Lambda_2^C, \Lambda_2^G\}, \dots, \{\Lambda_n^C, \Lambda_n^G\}\}$ 
3:  $\forall \tau_i \in \tau, \Lambda_i^C = \{\tau_{i,j} \in \tau_i, \Lambda_i^C = \emptyset \text{ and } E(k) = \emptyset\}$ 
4: for  $\tau_i \in \tau$  in descending order of priority do
5:   if  $\tau_i$  is unschedulable or  $\forall \tau_j \in \text{hp}(\tau_i)$  is schedulable with  $\Lambda$  by Eq. (9) then
6:     return  $\Lambda$ 
7:   end if
8:   for  $\tau_{i,j} \in \Lambda_i^C$  do
9:     Set  $\Lambda_{i,j}^C = \Lambda_i^C \setminus \{\tau_{i,j}\}$  and  $\Lambda_{i,j}^G = \Lambda_i^G \cup \{\tau_{i,j}\}$ 
10:    Set  $\Lambda = \{\Lambda_1^C, \Lambda_1^G\} \cup \Lambda \setminus \{\Lambda_i^C, \Lambda_i^G\}$ 
11:    if  $\tau_i$  and  $\forall \tau_j \in \text{hp}(\tau_i)$  is schedulable with  $\Lambda$  by Eq. (9) then
12:       $E(k) = E(k) \cup \{\tau_{i,j}\}$ 
13:    end if
14:   end for
15:   if  $E(k) \neq \emptyset$  then
16:     Find  $\tau_{i,k} \in E(k)$  according to Eq. (10)
17:      $\Lambda_i^C = \Lambda_i^C \cup \{\tau_{i,k}\}, \Lambda_i^G = \Lambda_i^G \setminus \{\tau_{i,k}\}$  and  $E(k) = \emptyset$ 
18:     go to Line 5
19:   end if
20: end for
    
```

Timing guarantee, but *accuracy loss*




Online GPU/CPU re-allocation

LaLaRAND: Flexible Layer-by-Layer CPU/GPU Scheduling for Real-Time DNN Tasks
IEEE **RTSS** 2021

Real-Time Scheduling

Schedule all real-time tasks
such that there is no single job deadline miss



Some contents are from “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo.

Real-Time Scheduling

■ Task models

- Single job
- Aperiodic
- Sporadic
- Periodic
- DAG-structured
- Gang-structured
- ...

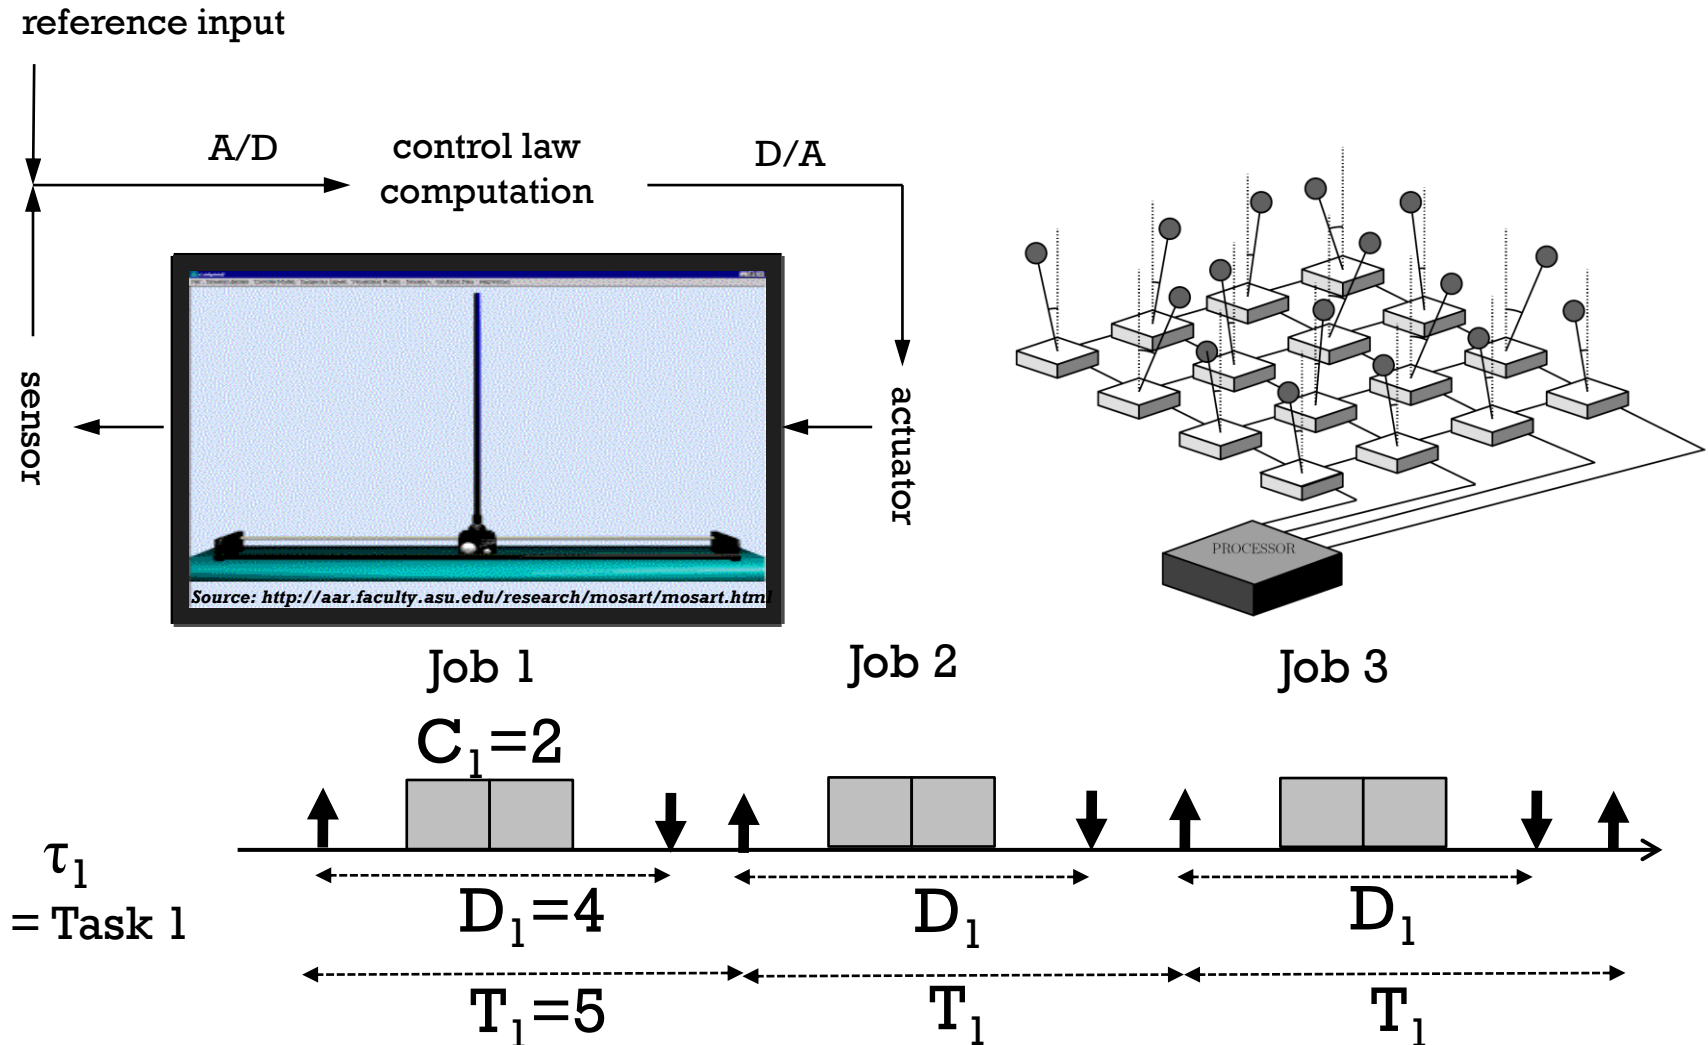
■ Constraints

- Precedence
- Shared resources
- Self-suspensions
- Mode changes
- Fully preemptive
- Fully non-preemptive
- Limited preemptive
- ...

■ Platforms

- Uniprocessor
 - ...
- Multiprocessor
 - ...
- GPU + CPU
- Distributed
- ...


Real-Time Scheduling



Real-Time Scheduling

■ Uniprocessor


$$T_1=4, D_1=4, C_1=1$$



Real-Time Scheduling

- Uniprocessor: at each time slot, only one job can be executed.

$$T_1=4, D_1=4, C_1=1$$




Schedule all real-time tasks
such that there is no single job deadline miss

= **Schedulable**


Real-Time Scheduling

- Uniprocessor: at each time slot, only one job can be executed.


$$T_1=4, D_1=4, C_1=1$$



$\tau_1 > \tau_2 > \tau_3$, non-preemptive



$\tau_1 > \tau_2 > \tau_3$, preemptive



Real-Time Scheduling

■ Task models

- Single job
- Aperiodic
- Sporadic
- **Periodic**
- DAG-structured
- Gang-structured
- ...

■ Constraints

- Precedence
- Shared resources
- Self-suspensions
- Mode changes
- **Fully preemptive**
- **Fully non-preemptive**
- Limited preemptive
- ...

■ Platforms

- **Uniprocessor**
 - ...
- Multiprocessor
 - ...
- GPU + CPU
- Distributed
- ...

Real-Time Scheduling

■ Optimization criteria

Schedule all real-time tasks
such that there is no single job deadline miss

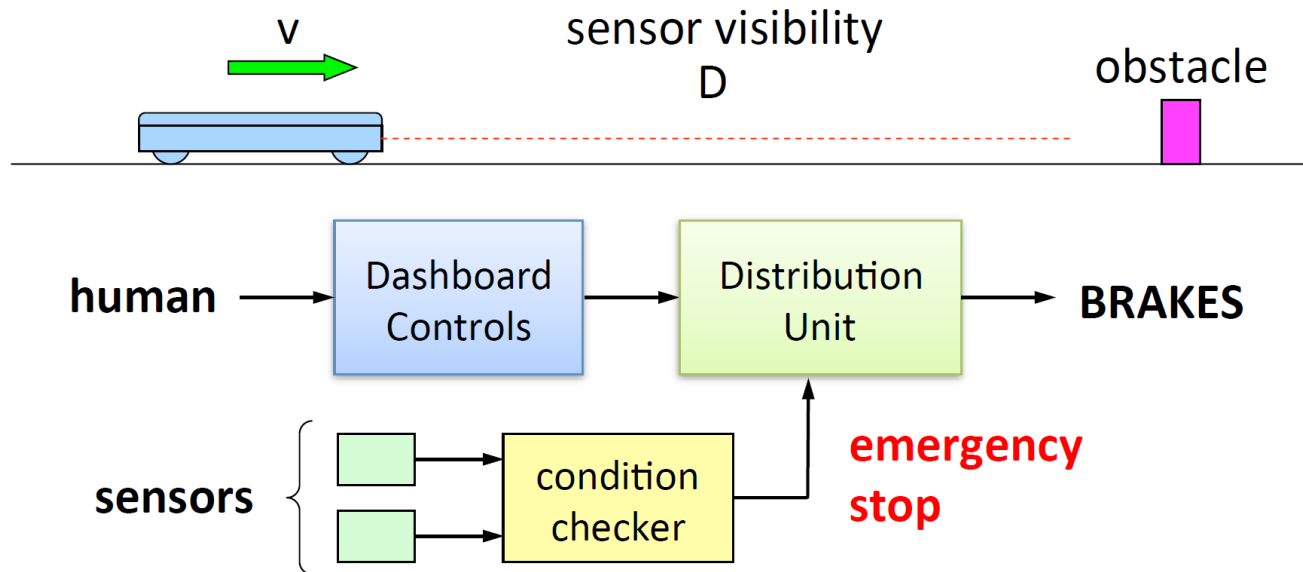
- Feasibility
- Response time
- Maximum lateness
- Utilization bandwidth
- Energy consumption
- Temperature
- Number of processors

Real-Time Scheduling

- Industrial/practical usage of real-time scheduling
 - Rate Monotonic is used in most industrial settings.
 - Priority Inheritance is in most RT kernels.
 - Sporadic Server is specified in POSIX.
 - CBS is implemented in LINUX.
 - EDF is now supported by a few kernels
 - Erika Enterprise: certified OSEK and adopted by Magneti Marelli in next generation ECUs
 - Ada 2005 runtime support
 - Linux: SCHED_DEADLINE in mainline since June 2014
 - A lot of RT Tools are now available for
 - WCET estimation, schedulability analysis, scheduling simulation, formal verification, etc.

Real-Time Scheduling

■ A practical example of real-time scheduling



GOAL: If an obstacle is detected, stop the train without hitting the obstacle.


PROBLEM: Find the sampling periods of the sensors that guarantee the feasibility of the goal

From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling


- Tasks are scheduled by Rate Monotonic (implicit deadlines)
- Let $\tau_s(C_s, T_s)$ be the periodic task devoted to sampling
- Assume τ_s has the shortest period (highest priority).
- Let U_{other} be the load of the other tasks




From "Real-Time Systems: Achievements and Perspectives" by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling



- Tasks are scheduled by **Rate Monotonic** (implicit deadlines)
- Let $\tau_s(C_s, T_s)$ be the periodic task devoted to sampling
- Assume τ_s has the **shortest period** (highest priority).
- Let U_{other} be the load of the other tasks



■ What is the physical meaning of $1/4 + 2/5 + 2/7$?

- Utilization
- What if the utilization is larger than 1.0?

From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling

The **minimum period** can be computed imposing the system schedulability by the Liu & Layland test:

The system is schedulable if $\frac{C_s}{T_s} + U_{other} \leq U_{lub}^{RM}$

that is $T_s \geq \frac{C_s}{U_{lub}^{RM} - U_{other}}$


$$T_s^{\min} = \frac{C_s}{U_{lub}^{RM} - U_{other}}$$

■ Example: $1/4 + 2/5 + 2/7 \leq 0.7$

From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling




From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling

The space covered by the train in $(T_s + \Delta + T_b)$ should not exceed



From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

- A practical example of real-time scheduling

$$v(T_s + \Delta) + X_b < D$$

$$X_b = \frac{v^2}{2\mu g}$$


$$v(T_s + \Delta) + \frac{v^2}{2\mu g} < D$$

$$T_s < \frac{D}{v} - \frac{v}{2\mu g} - \Delta \rightarrow T_s^{\max}$$

From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

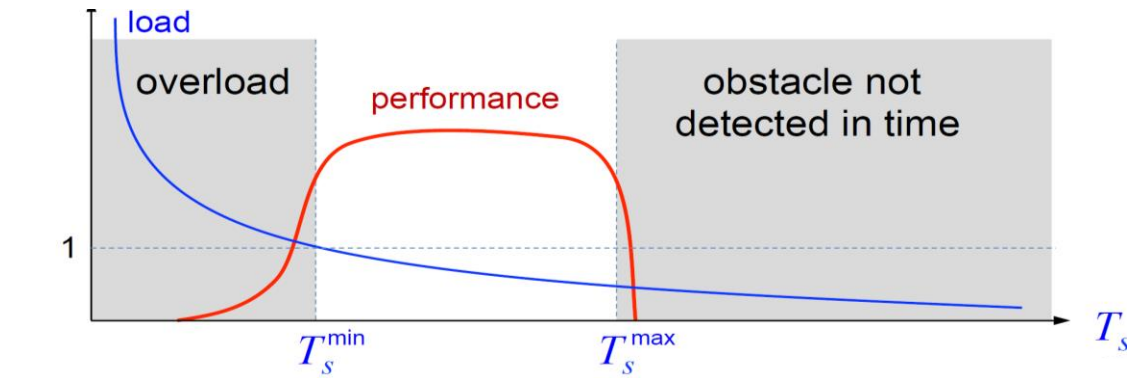
- A practical example of real-time scheduling



From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

■ A practical example of real-time scheduling



$$T_s^{\min} = \frac{C_s}{U_{\text{lub}}^{\text{RM}} - U_{\text{other}}}$$

$$T_s^{\max} = \frac{D}{v} - \frac{v}{2\mu g} - \Delta$$

Why this is possible?

The system is schedulable if $\frac{C_s}{T_s} + U_{\text{other}} \leq U_{\text{lub}}^{\text{RM}}$

Scheduling algorithm + Schedulability analysis

From “Real-Time Systems: Achievements and Perspectives” by Giorgio Buttazzo

Real-Time Scheduling

- Task model, constraints and platform are usually given.
- Then, how can we achieve timeliness of a real-time systems?

Schedule all real-time tasks
such that there is no single job deadline miss

■ Task models

- Single job
- Aperiodic
- Sporadic
- **Periodic**
- DAG-structured
- Gang-structured
- ...

■ Constraints

- Precedence
- Shared resources
- Self-suspensions
- Mode changes
- **Fully preemptive**
- **Fully non-preemptive**
- Limited preemptive
- ...

■ Platforms

- **Uniprocessor**
 - ...
- Multiprocessor
 - ...
- GPU + CPU
- Distributed
- ...

Real-Time Scheduling

■ From now on, we focus on the following *periodic/sporadic task model*

■ τ_i : Task i

■ T_i : the minimum separation / period

■ C_i : the worst-case execution time (WCET)


■ D_i : the relative deadline

■ Job: an instance of a task

(T_i, C_i, D_i)

Or

(T_i, C_i) , implying $T_i = D_i$

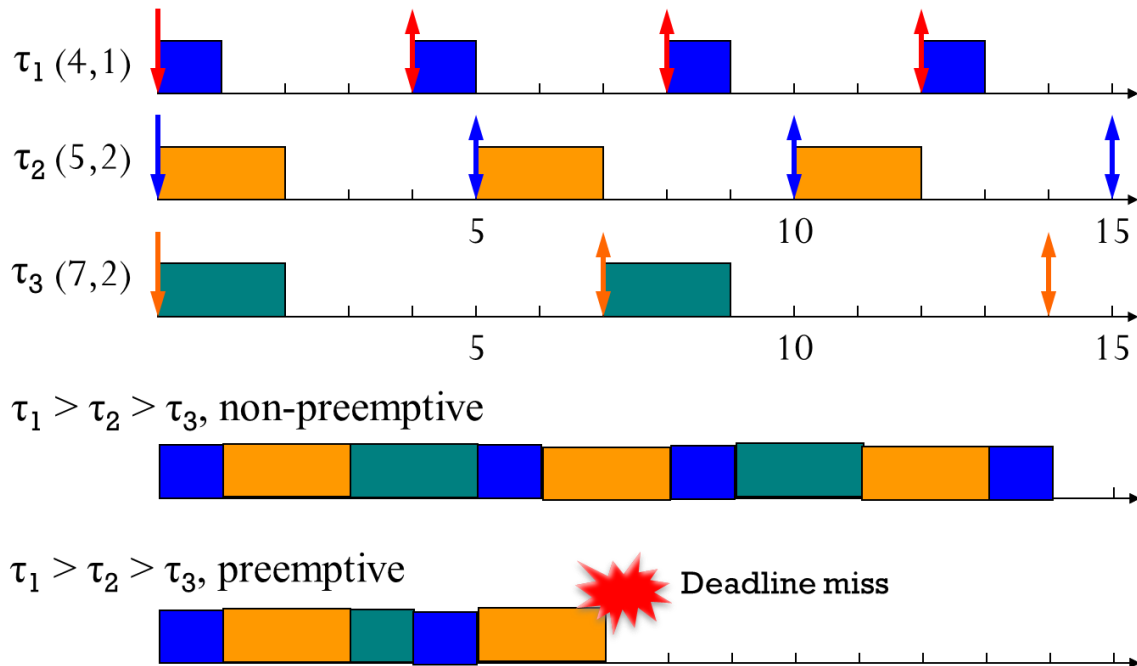


Real-Time Scheduling

- From now on, we focus on the following *periodic/sporadic task model*
 - τ_i : Task i
 - T_i : the minimum separation / period
 - C_i : the worst-case execution time (WCET)
 - D_i : the relative deadline
 - Where do these parameters come from?
 - T_i and D_i : usually from the target system (recall the hitting-obstacle example)
 - C_i : empirically measured, or code analysis
 - WCET analysis, under given platform

Real-Time Scheduling

- From now on, we focus on the *fully preemptive* constraint and a *uniprocessor platform*.
- A higher-priority job can preempt a currently-executing lower-priority job.
- At each time slot, only one job can be executed.




- Where do these constraints/platforms come from?
 - Usually from the target system and its design

Real-Time Scheduling


- Scheduling goal

Schedule all real-time tasks
such that there is no single job deadline miss

- Under the periodic/sporadic task model, the fully preemptive constraint, and a uniprocessor platform
- Two main components for real-time scheduling
 - Scheduling algorithm
 - Determine when each job executes (which job to be executed)
 - Schedulability analysis
 - Guarantee no deadline miss of a given task set under the target scheduling algorithm



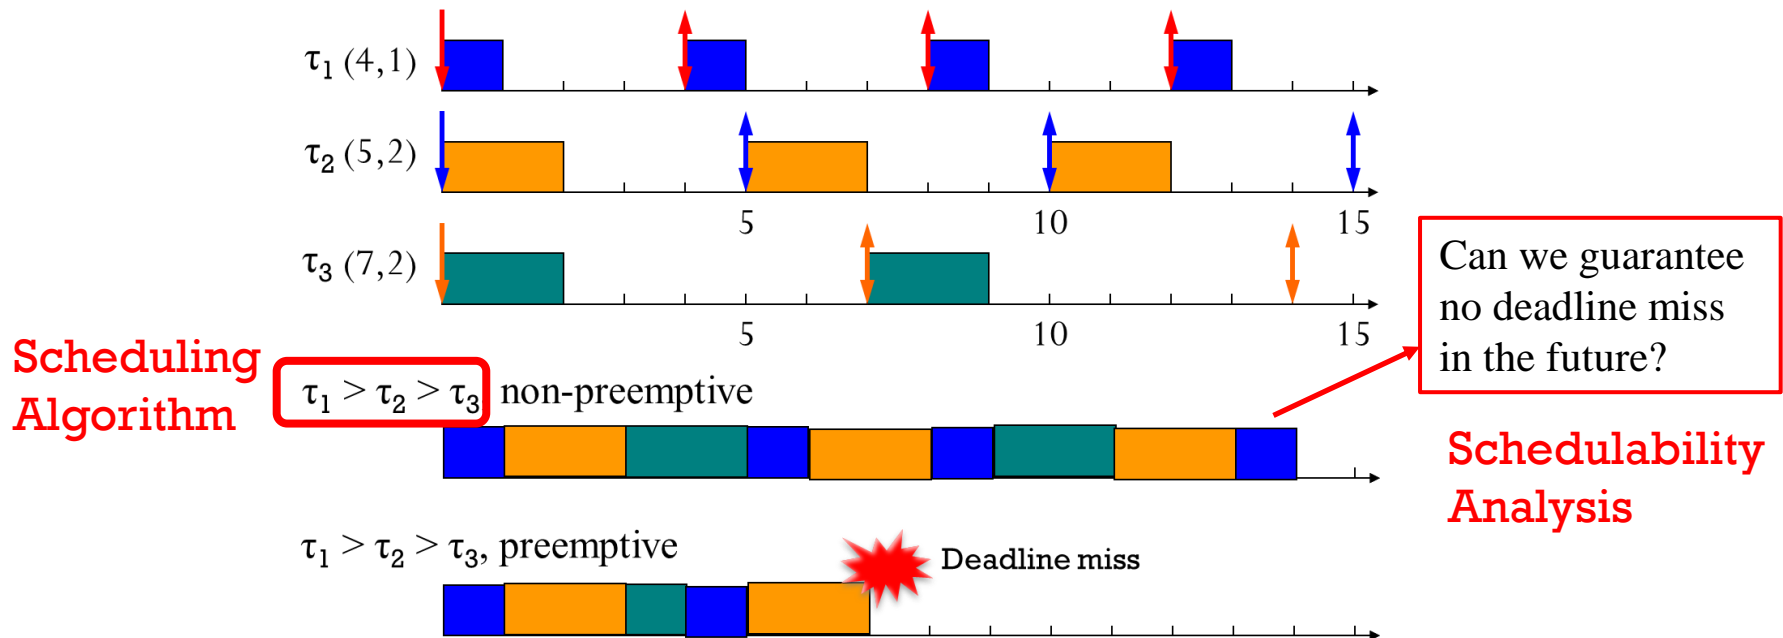
Exact (sufficient and necessary) 40



Only sufficient


Real-Time Scheduling

- Two main components for real-time scheduling
 - Scheduling algorithm
 - Determine when each job executes (which job to be executed)
 - Schedulability analysis
 - Guarantee no deadline miss of a given task set under the target scheduling algorithm



Real-Time Scheduling

- Why do we need schedulability analysis? Isn't it possible to guarantee schedulability simply by simulating a target set of tasks?
- Is it possible to simulate all different combinations of offsets?
- Can each simulation guarantee the schedulability even with actual execution times less than their corresponding WCETs?
- Is it feasible to simulate a target set of tasks if their LCM is very large?



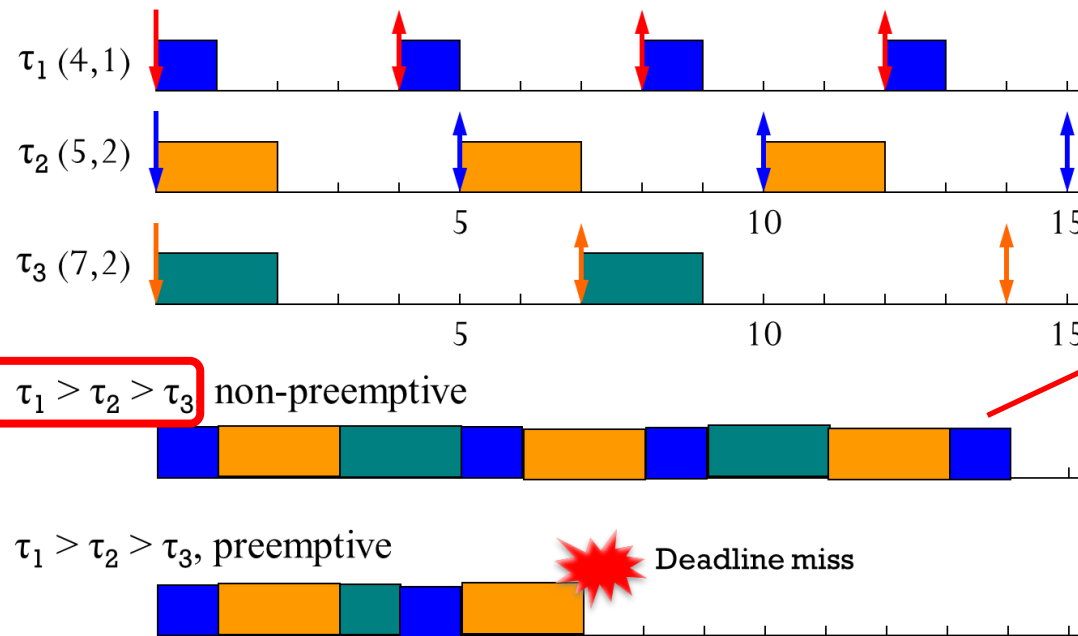
Real-Time Scheduling

- Which scheduling algorithm with which condition (i.e., schedulability analysis) guarantees that there is no single job deadline miss forever?
 - With any offset,
 - With any actual execution time,
 - With any arbitrary long period,

Example

The system is schedulable if $\frac{C_s}{T_s} + U_{other} \leq U_{lub}^{RM}$

Scheduling Algorithm




Can we guarantee no deadline miss in the future?

Schedulability Analysis

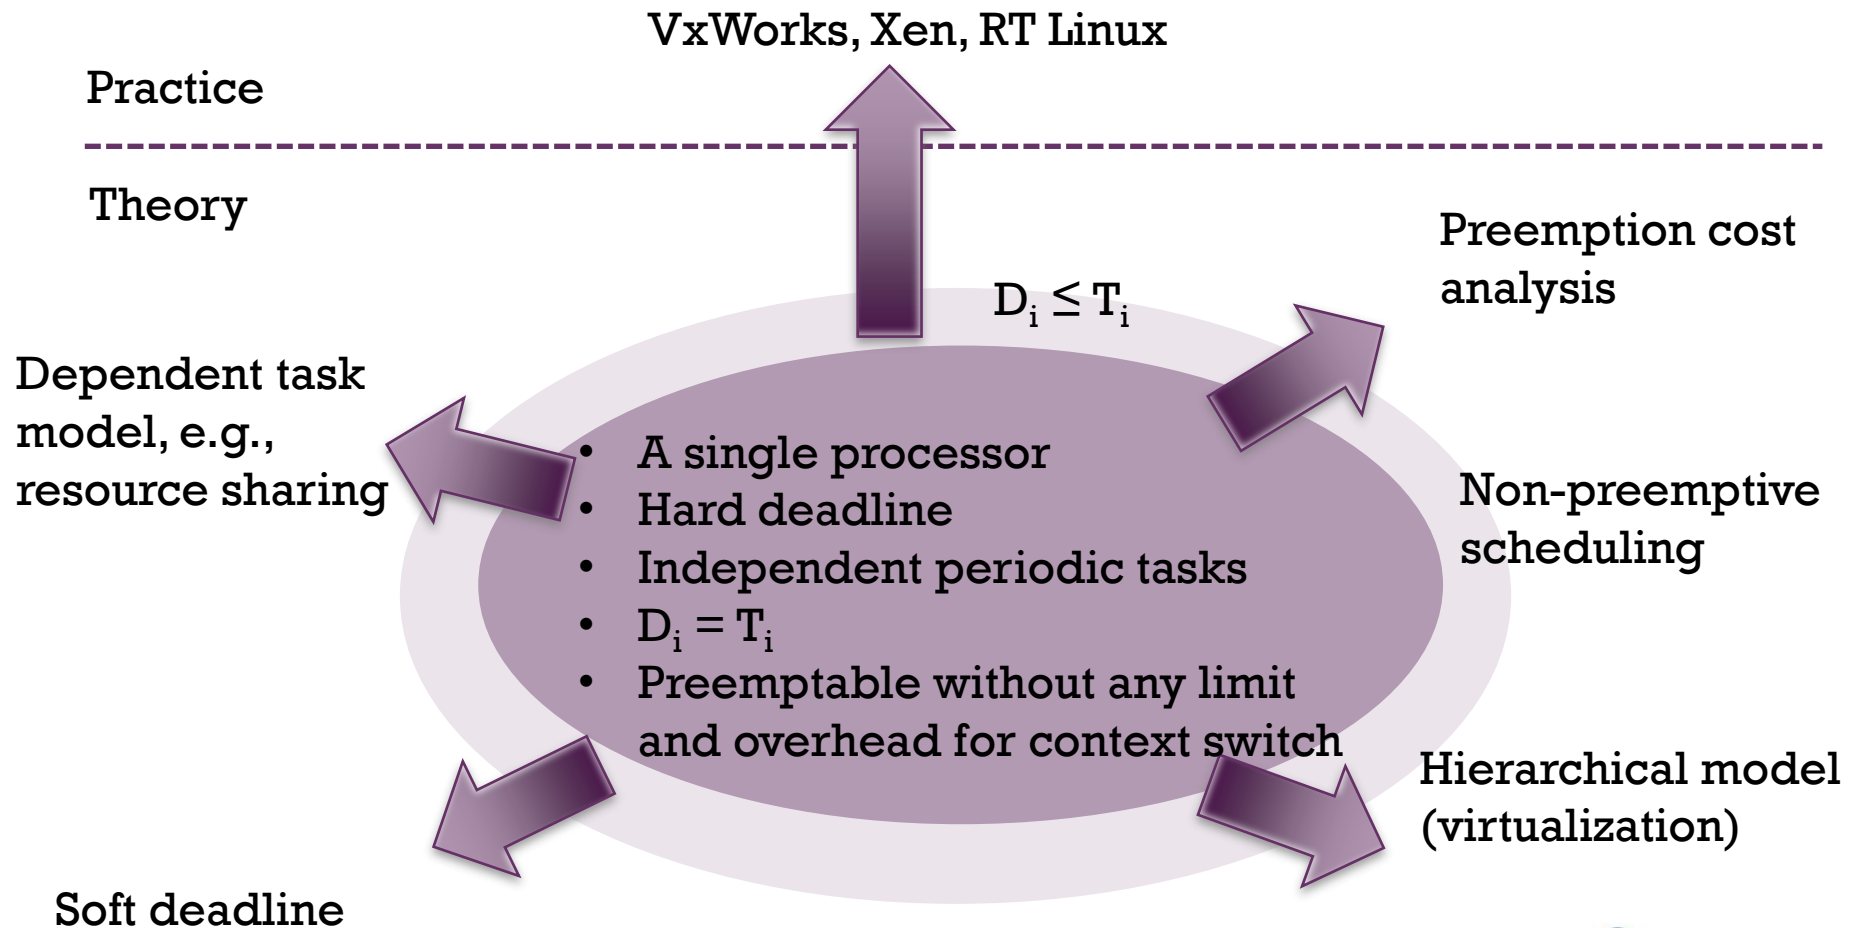
Real-Time Scheduling

- Let's focus on the most popular (simplest) setting.
 - A single processor
 - Hard deadline
 - Independent periodic tasks
 - Relative deadline = period ($D_i = T_i$)
 - Preemptable without any limit
 - No overhead for context switch



Real-Time Scheduling


- How can this setting be extended?



Clock-Driven Task Scheduling

■ Clock-driven


- a schedule determines (off-line) which job to be executed at each instant
- static or *cyclic*
- predictable and deterministic
- scheduler: invoked by a timer
- multiple *tables* for different *operation modes*



$$T_1 = 6, C_1 = 3, D_1 = 6$$

$$T_2 = 8, C_2 = 3, D_2 = 8$$

$$\text{Major cycle} = \text{lcm}(6, 8) = 24$$



Pros and Cons of Cyclic Schedule

■ Pros

- Simple, table-driven, easy to validate (knows what's going on at any moment)
- Fits well for *harmonic* periods and small system variations
- Static schedule => deterministic, static resource allocation, *no preemption*
- Small jitter


■ Cons

- Difficult to change (need to re-schedule all tasks)
- Fixed released times for the set of tasks
- Difficult to deal with different temporal dependencies
- Scheduling algorithm may become complex (NP-hard)
- Doesn't support aperiodic and sporadic tasks efficiently

Round-Robin Task Scheduling

■ Weighted Round-robin


- **Interleave** job executions
- Allocate a time slice to each job in the FIFO queue
- Time slice may **vary** while sharing the processor
- Good for pipelined jobs, e.g., network packets



Priority-Driven Task Scheduling

■ Priority-driven

- The highest-priority job gets to run until completion or blocked
- A processor is never idle if ready jobs are waiting (**work-conserving**)
- Preemptive or non-preemptive
- Priority assignment can be **static** or **dynamic**
- Scheduler just looks at the priority queue for waiting jobs (**list schedule**)



Priority-Driven Task Scheduling

- Why priority-driven scheduling algorithm?
 - Use priority to represent urgency/importance
 - Easy implementation of scheduler (compare task priorities and then dispatch tasks accordingly)
 - Tasks can be added or removed easily
 - No *direct control* of execution instant

Classification of Priority-Driven Scheduling

- **Task-level fixed-priority (TFP)**: all jobs of a periodic task have the same fixed priority

- RM (Rate-monotonic) [Liu and Layland 1973] — the higher the task frequency, the higher its priority



- **Task-level dynamic-priority**: different priorities for different jobs of a periodic task

- **Job-level fixed-priority (JFP)**: priority of each job is fixed

- EDF (Earliest Deadline First) [Liu and Layland 1973]

- **Job-level dynamic-priority (JDP)**: priority of each job is dynamic

- LLF (Least Laxity First) [Leung 1989]




Classification of Priority-Driven Scheduling

- **Task-level fixed-priority (TFP):** all jobs of a periodic task have the same fixed priority
 - RM (Rate-monotonic) [Liu and Layland 1973] — the higher the task frequency, the higher its priority


$\tau_1 (3, 1)$

Task₁
 $T_1=3$
 $C_1=1$



$\tau_2 (5, 3)$

Task₂
 $T_2=5$
 $C_2=3$




Classification of Priority-Driven Scheduling

- **Job-level fixed-priority (JFP):** priority of each job is fixed
 - EDF (Earliest Deadline First) [Liu and Layland 1973]


$\tau_1 (3, 1)$

Task₁
 $T_1=3$
 $C_1=1$




$\tau_2 (5, 3)$

Task₂
 $T_2=5$
 $C_2=3$



Classification of Priority-Driven Scheduling

- **Job-level dynamic-priority (JDP)**: priority of each job can change over time
 - LLF (Least Laxity First) [Leung 1989]



Classification of Priority-Driven Scheduling

- Can we guarantee no deadline miss of all jobs of a given task set under each scheduling algorithm?

- TFP: RM


- JFP: EDF

- JDP: LLF

- What is the maximum utilization?

- $U = \sum C_i/T_i \leq 1.0$

- Example: $1/4 + 2/5 + 2/7$



- Can each scheduling algorithm achieve the maximum utilization?

RM: Overview

- Critical instant under RM

- Utilization-based analysis

- $U \leq n(2^{1/n} - 1.0)$

- $n \rightarrow \infty, n(2^{1/n} - 1.0) \rightarrow 0.69$

- Only sufficient

Without loss of generality

$$T_1 \leq T_2 \leq \dots \leq T_n$$

- Response-time analysis




- Exact

- Some applications

- Schedulability with transient overhead, interrupts and blocking

RM: Critical Instant

- Critical instant of τ_i
 - A job of τ_i arriving at that instant has the largest response time
- Which τ_2 has the maximum response time?




- In-phase instant is critical under RM
 - All higher priority tasks are released at the same instant of $J_{i,c}$ (assume all jobs are completed before the next job is released)

Why?

RM: Utilization-Based Analysis


Proof of $U \leq n(2^{1/n} - 1.0)$

- Let's consider two tasks
 - Task₂ can only be executed when Task₁ is not in the system.
- Let $T_2 < 2T_1$. Determine the maximum schedulable C_2 .
 - If $T_2 \leq T_1 + C_1$, $\max(C_2) = T_1 - C_1$
 - $U = C_1/T_1 + (T_1 - C_1)/T_2$
 - Else, $\max(C_2) = T_2 - 2C_1$
 - $U = C_1/T_1 + (T_2 - 2C_1)/T_2$



RM: Utilization-Based Analysis

- Given C_1 , T_1 , and T_2 , plot U
- The minimum U occurs when
 - $T_2 = T_1 + C_1$where $U = C_1/T_1 + (T_1 - C_1)/(T_1 + C_1)$



- What is the minimum U ?
 - Take the derivative w.r.t T_1 and set $dU/dT_1 = 0$
 - We will get $C_1 = (2^{1/2} - 1.0)T_1$ and $U = 0.828$
- General cases with n
 - $$-C_1 * T_1^{-2} - (T_1 - C_1) * (T_1 + C_1)^{-2} + (T_1 + C_1)^{-1} = 0$$
 - $$-C_1 * T_1^{-2} + 2C_1 * (T_1 + C_1)^{-2} = 0$$

RM-schedulable if $U \leq n(2^{1/n} - 1.0)$

RM: Utilization-Based Analysis

■ Example

	C_i	T_i	U_i
Task ₁	20	100	0.200
Task ₂	40	150	0.267
Task ₃	100	350	0.286

■ Total utilization is

$$.200 + .267 + .286 = .753 < U(3) = 3(2^{1/3} - 1.0) = .779$$

■ The periodic tasks in the example are schedulable according to the UB test.

RM: Utilization-Based Analysis

■ Example

	C_i	T_i	U_i
Task ₁	40	100	0.400
Task ₂	40	150	0.267
Task ₃	100	350	0.286

■ Total utilization is

$$.400 + .267 + .286 = .953 > U(3) = 3(2^{1/3} - 1.0) = .779$$

■ UB test cannot deem the task schedulable.

■ Try the exact analysis

RM: Response-Time Analysis

- Response time analysis

- The duration between a job's release time and finishing time

Without loss of generality

$$T_1 \leq T_2 \leq \dots \leq T_n$$

$$a_{n+1} = C_i + \left(\sum_{j=1}^{i-1} \left\lceil \frac{a_n}{T_j} \right\rceil C_j \right) \quad \text{where} \quad a_0 = \sum_{j=1}^i C_j$$

A set of higher-priority tasks

- Test terminates when $a_{n+1} = a_n$
- Task i is schedulable if its response time is before its deadline: $a_n \leq T_i$
 - a_n is the response time of Task $_i$
- Can be applied to any TFP
 - *Not only RM, but also any TFP*


RM: Response-Time Analysis

Without loss of generality

$$T_1 \leq T_2 \leq \dots \leq T_n$$


$$a_{n+1} = C_i + \left[\sum_{j=1}^{i-1} \left\lceil \frac{a_n}{T_j} \right\rceil C_j \right] \quad \text{where} \quad a_0 = \sum_{j=1}^i C_j$$

A set of higher-priority tasks



RM: Response-Time Analysis

- How does this work? Task₃ is schedulable?



	C_i	T_i	U_i
Task ₁	40	100	0.400
Task ₂	40	150	0.267
Task ₃	100	350	0.286

$$a_0 = \sum_{j=1}^3 C_j = C_1 + C_2 + C_3 = 40 + 40 + 100 = 180$$

$$\begin{aligned}
 a_1 &= C_i + \sum_{j=1}^{i-1} \left\lceil \frac{a_0}{T_j} \right\rceil C_j = C_3 + \sum_{j=1}^2 \left\lceil \frac{a_0}{T_j} \right\rceil C_j \\
 &= 100 + \left\lceil \frac{180}{100} \right\rceil (40) + \left\lceil \frac{180}{150} \right\rceil (40) = 100 + 80 + 80 = 260
 \end{aligned}$$

RM: Response-Time Analysis

$$a_2 = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{a_1}{T_j} \right\rceil C_j = 100 + \left\lceil \frac{260}{100} \right\rceil (40) + \left\lceil \frac{260}{150} \right\rceil (40) = 300$$

$$a_3 = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{a_2}{T_j} \right\rceil C_j = 100 + \left\lceil \frac{300}{100} \right\rceil (40) + \left\lceil \frac{300}{150} \right\rceil (40) = 300$$

$$a_3 = a_2 = 300$$

Done!

■ Task is schedulable using RT test.

■ $a_3 = 300 < T_3 = 350$


RM: Response-Time Analysis

■ Why does this work?

- Time-demand analysis based on critical instant
- If J_i is done at t , then the total work must be done in $[0, t]$ is (from J_i and all higher priority tasks): **time demand function**

$$w_i(t) = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{T_k} \right\rceil C_k$$

- Can we find $t \leq D_i$ such that $w_i(t) \leq t$
 - Cannot check all $t \in [0, d_i]$
 - Check all **arrival instants** and d_i



RM: Transient Overload

- Question: what if the task with a smaller period is not important to the underlying application?
- Answer: consider period transformation, period aggregation or period splitting
- Example: consider the following unschedulable task set

τ_i	C_i	EC_i (avg exec time)	T_i	comments
τ_1	20	10	100	
τ_2	30	25	150	
τ_3	80	40	210	Non-critical
τ_4	100	20	400	

RM: Transient Overload

- **Solution 1:** reduce Task₃'s priority by lengthening its period, possible only if Task₃'s relative deadline can be greater than its original period. In such a case, replace Task₃ by two tasks Task₃⁰ and Task₃⁰⁰, each with period 420, WC exec times $C_3^0 = C_3^{00} = 80$, avg exec times $a_3^0 = a_3^{00} = 40$. Task₃⁰ and Task₃⁰⁰ must be phased to be 210 time units apart. If the set {Task₁, Task₂, Task₃⁰, Task₃⁰⁰, Task₄} is RM-schedulable, done.

τ_i	C_i	EC_i (avg exec time)	T_i	comments
τ_1	20	10	100	
τ_2	30	25	150	
τ_3	80	40	210	Non-critical
τ_4	100	20	400	

RM: Transient overload

- **Solution 2:** increase Task₄'s priority by splitting each invocation into two: Task₄⁰: $C_4^0 = C_4/2$, $a_4^0 = a_4/2$, $T_4^0 = T_4/2$
 - {Task₁, Task₂, Task₄} or {Task₁, Task₂, Task₄⁰} are schedulable

τ_i	C_i	EC_i (avg exec time)	T_i	comments
τ_1	20	10	100	
τ_2	30	25	150	
τ_3	80	40	210	Non-critical
τ_4	100	20	400	

EDF: Schedulability Analysis


■ RM's critical instants hold for EDF?

■ Why or why not?

■ Critical instant of τ_i


- A job of τ_i arriving at that instant has the largest response time

■ Which τ_2 has the maximum response time?



■ In-phase instant is critical under RM

- All higher priority tasks are released at the same instant of $J_{i,c}$ (assume all jobs are completed before the next job is released)




EDF: Schedulability Analysis

■ EDF is schedulable iff $U \leq 1.0$

■ $U = 1/4 + 2/5 + 2/7$

■ How to prove?

■ Using busy interval



EDF: Schedulability Analysis

■ Example

■ A digital robot with EDF schedule

■ Control loop: $C_c \leq 8\text{ms}$ at 100Hz

■ BIST (Built-in-Self-Testing): $C_b \leq 50\text{ms}$

■ BIST can be done every 250ms

$$u_c + u_b = \frac{8}{10} + \frac{50}{T_b} \leq 1$$

■ Add a telemetry task to send and receive messages with $C_t \leq 15\text{ms}$

■ If BIST is done every 1000ms

■ The telemetry task can have a relative deadline of 100ms

■ Sending or receiving must be separated by at least 100ms

$$u_c + u_b + u_t = \frac{8}{10} + \frac{50}{1000} + \frac{15}{D_t} \leq 1$$

Summary: $D_i = T_i$

- TFP: RM
 - $U \leq n(2^{1/n} - 1.0)$
 - Response-time analysis (exact)
- JFP: EDF
 - $U \leq 1$ (exact)
- JDP: LLF
 - $U \leq 1$ (exact)
- The three algorithms are **optimal** for each category.
- All the results hold for sporadic tasks.

Extension: $D_i = T_i \rightarrow D_i \leq T_i$


- TFP: DM (Deadline Monotonic)
 - $\sum C_i/D_i \leq n(2^{1/n} - 1.0)$
 - Response-time analysis: $a_n \leq D_i$, still exact
- JFP: EDF
 - $\sum C_i/D_i \leq 1$
 - Not exact, but there exists another exact analysis.
- JDP: LLF
 - $\sum C_i/D_i \leq 1$
- The three algorithms are still **optimal** for each category.
- All the results hold for sporadic tasks.

Extension: $D_i = T_i \rightarrow D_i \leq T_i$


■ JFP: EDF

■ Any exact analysis?

$$\text{DBF}_{\text{EDF}}(\tau_i, l) \triangleq \left(\left\lfloor \frac{l - D_i}{T_i} \right\rfloor + 1 \right) \cdot C_i.$$



(a) $\text{DBF}_{\text{EDF}}(\tau_i, l)$



Extension: $D_i = T_i \rightarrow D_i \leq T_i$

■ JFP: EDF

■ Any exact analysis?

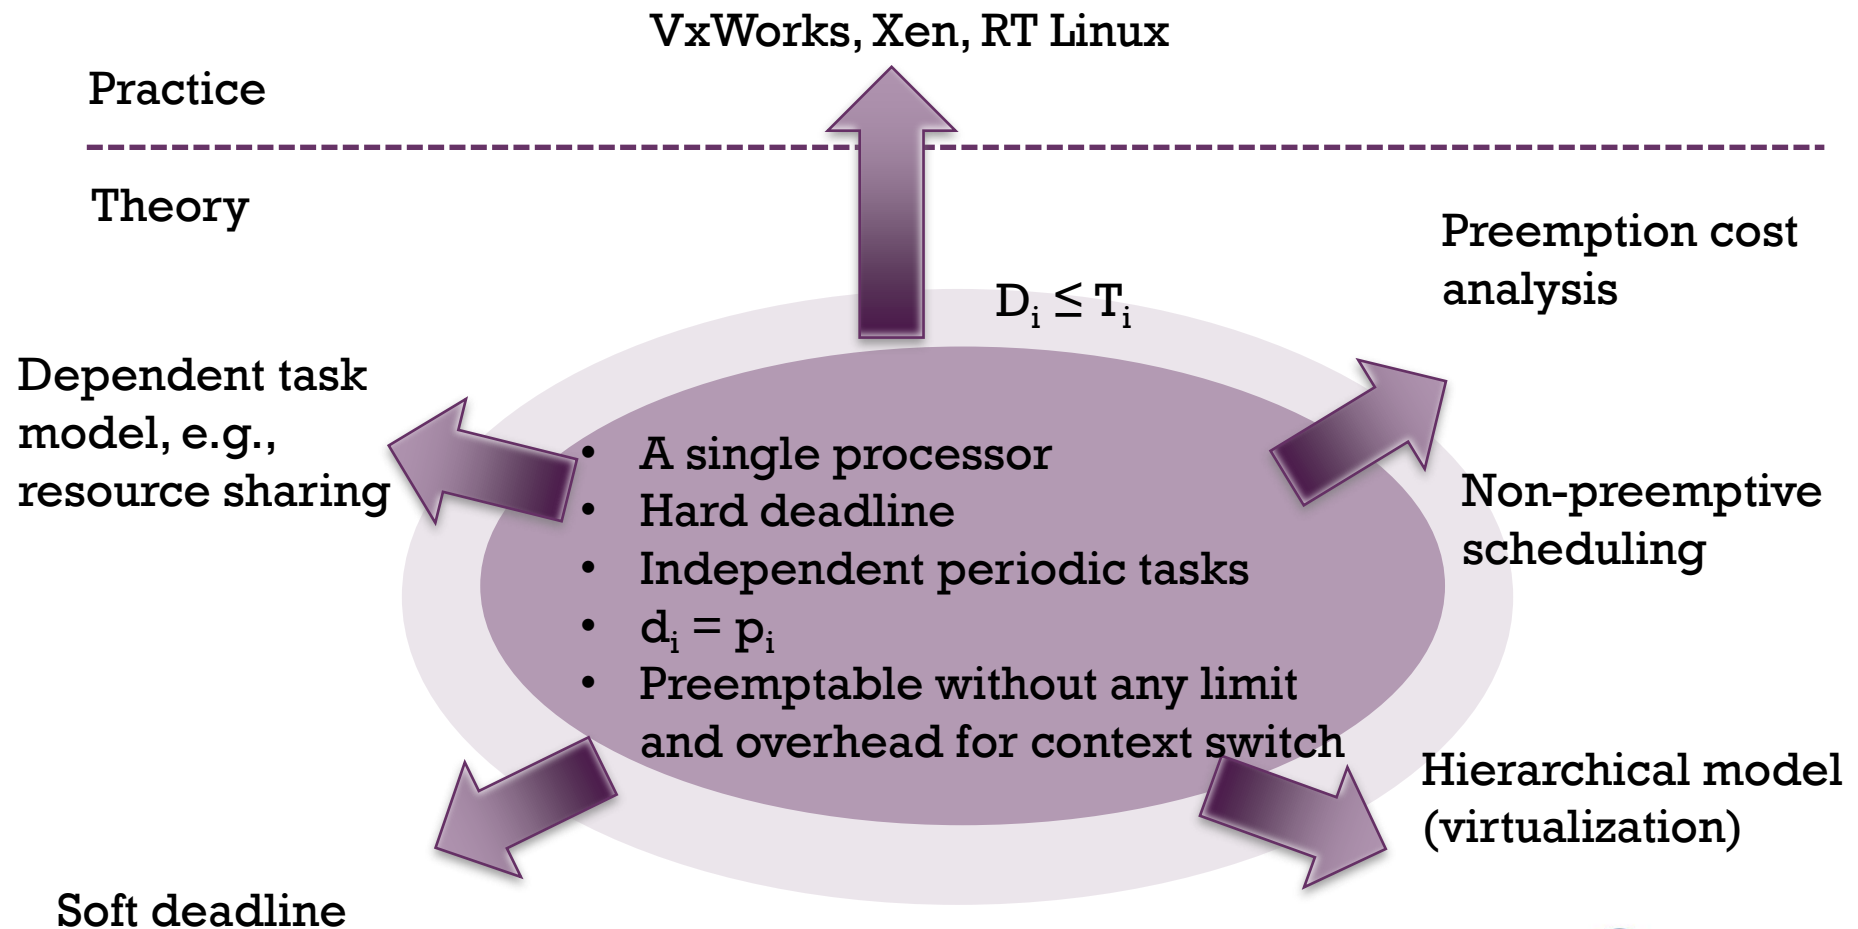
- $\sum \text{DBF}(\text{Task}_i, t) \leq t$, for all $t > 0$

$$\text{DBF}_{\text{EDF}}(\tau_i, l) \triangleq \left(\left\lfloor \frac{l - D_i}{T_i} \right\rfloor + 1 \right) \cdot C_i.$$

■ Proof

- Busy interval

Where are we?



How to Apply This to a Target System?

- Check

- Task model

- Any dependency?
 - Any shared resources?
 - (T, C, D) is given?

- Platform

- Constraints